

# Fortran OP2 User's Manual

Carlo Bertolli, Adam Betts  
Imperial College London, UK

May 19, 2011

# Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Overview</b>	<b>2</b>
2.1	Sets . . . . .	2
2.2	Maps . . . . .	2
2.3	Datasets . . . . .	2
2.4	Global Variables . . . . .	3
2.5	Parallel Computation over Sets . . . . .	3
<b>3</b>	<b>OP2 Data Types</b>	<b>4</b>
<b>4</b>	<b>Initialisation and Termination Routines</b>	<b>4</b>
<b>5</b>	<b>Data Structure Declaration Routines</b>	<b>5</b>
<b>6</b>	<b>Parallel Computation Routines</b>	<b>6</b>
<b>7</b>	<b>User Kernel Interface</b>	<b>7</b>

## 1 Introduction

OP2 is a library targeting parallel Computation Fluid Dynamic (CFD) applications. It currently supports the Finite Volume Method and research work is in progress to extend it to support the Finite Element Method. OP2 supports both structured and unstructured meshes, but it is especially optimised to the unstructured case. The OP2 compiler targets CUDA, OpenCL and AVX/SSE as backends, i.e. it transforms an OP2 program in a semantically corresponding program implemented in the user-selected backend. This documents describes the abstractions provided by the OP2 library and it includes its Fortran API.

## 2 Overview

In this section we give an overview of the abstractions provided by OP2.

### 2.1 Sets

An OP2 set is used to model the basic entities forming a mesh. Examples of sets are vertices, edges, etc. In OP2, the only application-related information associated to a set is its size, i.e. the number of elements it includes. In a typical situation, this is an information that is contained in user-defined input files. The actual topology of a mesh is defined in terms of relations between OP2 sets, i.e. maps (see next Subsection).

### 2.2 Maps

An OP2 map relates two sets  $A$  and  $B$ , where each element  $a \in A$  ( $A$  is denoted *source set*) corresponds to one or more elements in the set  $B$  (denoted *target set*), i.e.  $a \rightarrow \{b_1, \dots, b_N\} \in B$ , where  $N$  is denoted as the *dimension* of the map and it must be the same constant value for all elements in  $A$ . For instance, in a typical situation an element in the edges set corresponds to two elements in the vertices set. This also means that it is not possible to define maps with variable dimension (i.e. a different dimension depending on the source set element), an example of which is the mapping between vertices and edges in a mesh.

Also the mapping information is typically stored in a user-defined input file, which data is used to populate arrays that are passed to the routine initialising a map.

### 2.3 Datasets

An OP2 dataset is associated to a specific OP2 set, where we relate a set of  $K$  scalar values  $d_1, \dots, d_K$  to each element of the set. The value  $K$  is denoted as the *dimension* of the dataset and it must be the same constant value for all set elements.

Dataset values are typically obtained from an user-defined input file, and the number of scalar values must be equal to the corresponding set size multiplied by the dataset dimension (i.e. to correctly assign the same amount of data to each set element).

## 2.4 Global Variables

An OP2 global variable is not associated to a specific *op\_set*, but it is shared by all parallel loop units<sup>1</sup>. Each parallel loop unit can access all of its fields, without the need of a map.

## 2.5 Parallel Computation over Sets

The subroutine used to express a parallel computation in OP2 is called *op\_par\_loop*. An *op\_par\_loop* is defined over a specific *op\_set*, called the *target* set. The effect of an *op\_par\_loop* is to read and update *op\_dat* variables, according to a user-specified subroutine. In general terms, for each parallel computation the programmer has to specify:

- The target *op\_set*, over which the parallelism is defined. Semantically, the *op\_par\_loop* instantiate a sequential computation for each element in the *op\_set* over which it is defined.
- A “user kernel” subroutine, to be evaluated once for each element in the *op\_set* over which the *op\_par\_loop* is defined. We will always refer to these user-defined, application-level routines as “user kernels”, to distinguish them from the notion of a CUDA kernel, which is one of the target programming models of OP2. The general form of a user kernel subroutine is described in Section 7.
- A list of *op\_dat* variables that are read and modified by the user kernel subroutine. Each of these variables is also a parameter to the user kernel. It is task of the OP2 run-time to provide each user kernel invocation with the proper section of input *op\_dat* arguments. Notice that, given a *op\_set*  $S$  over which an *op\_par\_loop* is defined, eligible *op\_dat* variables for the parallel loop are all those that are defined for  $S$ , but also all those that are defined for another *op\_set*  $T$  if the user has declared an *op\_map* between  $S$  and  $T$ . The former case, in which the passed *op\_dat* is defined over  $S$  is said to be a *direct* case, as the *op\_dat* can be accessed directly given an index of an element in  $S$ . The latter case, in which the passed *op\_dat* is defined over  $T$ , it is said to be an *indirect* case, as the *op\_dat* elements can be accessed by: first transforming an element index of  $S$  into  $T$  (i.e. selecting one of the associated index if the map dimension is greater than 1); then using the transformed index in  $T$  to access the *op\_dat*.

To express the above behaviour, for each *op\_dat* passed to a *op\_par\_loop* the programmer needs to specify:

- the *op\_dat* variable;
- an index that: (i) in case of a directly accessed *op\_dat* or a global data, it is not used and it must be set to *OP\_NONE*; (ii) in case of indirectly accessed *op\_dat*, via a given map  $M$  (e.g. from set  $S$  to  $T$ ) an integer index in the range  $[1, \text{dim}]$ , where  $\text{dim}$  is the dimension of  $M$ . In this case, the index identifies which of the mapped elements in  $T$  is to be considered;
- an *op\_map* variable: in case of directly accessed *op\_dat*, the programmer must specify the keyword *OP\_ID*. In case of a global data the programmer must specify the keyword *OP\_GBL*. Finally, in case of an indirectly accessed *op\_dat*, the user must specify the required *op\_map* variable needed to access the *op\_dat*.
- an accessing operation code, which must be one of the following named constants: *OP\_READ*, *OP\_WRITE*, *OP\_RW*, and *OP\_INC*. These accessing modes must be followed in the implementation of the user kernel. The *OP\_RW* operation denotes that the corresponding *op\_dat* variable will be both read and updated in the user kernel. The *OP\_INC* operation denotes that the related *op\_dat* elements will be incremented. In case of *OP\_INC* for a global variable, the effect is a reduction of the values computed by each kernel invocation.

Semantically, the OP2 run-time will instantiate a number of *parallel loop units* equal to the size of the iteration set. Each parallel loop unit, historically denoted with the term Virtual Processor (or VP, in short) in the parallel computing terminology, has the task of applying the user-kernel subroutine for a specific set element, by properly accessing the input datasets with their corresponding set element index.

<sup>1</sup>See below for the definition of parallel loop unit.

### 3 OP2 Data Types

#### op\_set

**Description** : this data type is used to express sets defining a mesh. OP2 require users to specify, for each set, its size, corresponding to the number of elements it is formed of, and a string of characters indicating its name. Figure 1 shows the declaration of multiple op\_set variables. As it can be noticed from the example, op\_set is implemented as a Fortran derived type (see *type* Fortran keyword used in the example).

```
type ( op_set ) :: vertices , edges , bedges , cells
```

Figure 1: Example of declaration of op\_set variables.

#### op\_map

**Description** : this data type is used to express a relation between two OP2 sets, defining the actual topology of an unstructured mesh, i.e. how its elements are related/connected to each other. If  $A$  and  $B$  are two op\_set variables, a map from  $A$  to  $B$  associates to each element  $a \in A$  a subset of elements  $b_1, \dots, b_n \in B$ . The number  $n$  must be equal for all elements  $a \in A$ , i.e. it is not possible that two elements in  $A$  are associated to a different number of elements in  $B$ . The number  $n$  is said to be the *dimension* of the op\_map, and it must be declared when initialising an op\_map as a constant value (see *\_map* call). Figure 2 shows the declaration of multiple op\_map variables, again as Fortran structures.

```
type ( op_map ) :: verts2Edges , edges2Verts , cells2Verts , cells2Bedges
```

Figure 2: Example of declaration of op\_map variables.

#### op\_dat

**Description** : this data type is used to express application-related data associated to a OP2 set, i.e. it is used to associate to each element of an op\_set a fixed number of data items. The number of data items associated to each op\_set element is denoted the op\_dat *dimension* and it must be a constant value (see *op\_decl\_dat* call). An op\_dat variable can encapsulate any primitive and structured Fortran type (i.e. defined with the *type* keyword). Figure 3 shows the declaration of multiple op\_dat variables, as Fortran derived types.

```
type ( op_dat ) :: p-q , p-qold , bound
```

Figure 3: Example of declaration of op\_dat variables.

## 4 Initialisation and Termination Routines

#### op\_init ( debug\_level )

**Description** : this routine initialises the OP2 environment and it must be called before any other OP2 subroutine.

**Dummy arguments** :

- integer(4) :: **debug\_level**: required debug level. It can assume the following values:
  - 0: none
  - 1: error checking

- 2: information on plan construction
- 3: report execution of parallel loops
- 4: report use of old plans
- 5: report positive checks in `op_plan_check` function

**op\_exit()**

**Description** : this routine terminates the OP2 environment and must be called after all OP2 subroutines.

## 5 Data Structure Declaration Routines

**op\_decl\_set ( size, setVariable, setName )**

**Description** : this routine is used to initialise a `op_set` variable. Its parameters are:

- `integer(4) :: size`: the number of elements in the `op_set`.
- `type(op_set) :: setVariable`: the `op_set` variable to be initialised.
- `character(len=*) :: setName`: a Fortran string specifying the name of the set (for debug purposes). An example of how a set name can be declared is shown in Figure 4.

<pre><b>character(len=5) :: nodesName = 'nodes'</b></pre>
---

Figure 4: Example of declaration of an `op_set` name, to be passed to the `op_decl_set` routine.

**op\_decl\_map ( sourceSet, targetSet, dimension, inputData, mapVariable, mapName )**

**Description** : this routine is used to initialise an `op_map` variable. According to the description above, an `op_map` relates the elements in a `op_set` (called *source set*) to the elements of another `op_set` (called *target set*). The parameters of this call are:

- `type(op_set) :: sourceSet`: the variable representing the source set of the declared map.
- `type(op_set) :: targetSet`: the variable representing the target set of the declared map.
- `integer(4) :: dimension`: the dimension of the map (i.e. the number of elements in `toSet` associated to each element in `fromSet`).
- `integer(4), dimension(*) :: inputData`: an array of integers, including the data needed to instantiate the mapping between `fromSet` and `toSet`. The length of this array must be equal to the size of the `fromSet` multiplied by the dimension of the map (see previous parameter).
- `type(op_map) :: mapVariable`: the `op_map` variable to be initialised.
- `character(len=*) :: mapName`: a Fortran string specifying the name of the map (for debug purposes).

**op\_decl\_dat ( associateSet, dimension, inputData, datVariable, datName )**

**Description** : this routine is used to initialise an `op_dat` variable. Its parameters are:

- `type(op_set) :: associatedSet`: the `op_set` variable to which the declared `op_dat` variable is associated.
- `integer(4) :: dimension`: the dimension of the declared `op_dat`, i.e. the number of data items associated to each element in `op_set`.
- `“type”, dimension(:) :: inputData`: an array of arbitrary type (here denoted with the “type” string, which can be `integer(4)`, `real(4)`, `real(8)`, etc..) including the actual data to be associated to the `op_set`. The length of this array must be equal to the size of the associated set multiplied by the dimension of the declared `op_dat` (see previous parameter). Therefore, the `:` character must be replaced by the actual array size.
- `character(len=*) :: mapName`: a Fortran string specifying the name of the dataset (for debug purposes).

**op\_decl\_gbl ( inputData, dimension, datVariable, gblName )**

**Description** : this routine is used to initialise a global OP2 data variable. This kind of data is represented, in terms of data types, as an op\_dat, but semantically with a special accessing policy (see Overview Section). The parameters of this call are:

- “type”, dimension(:) :: *inputData*: an array of arbitrary type including the actual data to be associated to the declared global variable. Currently, also scalar data must be expressed in terms of a single-sized array.
- integer(4) :: *dimension*: the size of the input data array, which is denoted as the dimension of the global data.
- type(op\_dat) :: *datVariable* the op\_dat variable to which the declared global variable is associated.
- character(len=\*) :: *mapName*: a Fortran string specifying the name of the global data (for debug purposes).

## 6 Parallel Computation Routines

The prototype for an op\_par\_loop call with two op\_dat input variables is shown below:

```
op_par_loop ( userSubroutine , iterationSet ,    &
              & datVar1 , index1 , map1 , access1 , &
              & datVar2 , index2 , map2 , access2 &
              & )
```

**Description** : as described in the Overview section, this routine is used to express a parallel computation over a specific op\_set, reading and modifying passed op\_dat variables. The op\_par\_loop call can take an arbitrary number of op\_dat variables (i.e. an instance of the above specified parameter list), which must correspond to the number of parameters to the user kernel passed as first argument (see below). The first two parameters of this call are:

- subroutine *userSubroutine*: this is the user kernel subroutine, that is instantiated in a number of parallel units equal to the size of the passed op\_set (see next parameter).
- type(op\_set) :: *iterationSet*: this is the set over which the op\_par\_loop is defined, and its size defines the number of parallel units that will be instantiated by the OP2 run-time when executing this call.

These parameters are followed by an arbitrary number of “lines”, each related to an input dataset. For each line, the programmer must specify:

- type(op\_dat) :: *datVar1*: the related dataset variable.
- integer(4) :: *index1*: in case of direct mapping or global data, this parameter is not used, and it must be set to the named constant OP\_NONE. Consider the case of indirectly accessed op\_dat, with a map from the iteration set  $S$  to another set  $T$  with dimension  $N$ . *index1* is a value in the range  $[1, N]$ , identifying one of the  $N$  elements in  $T$  to which each  $S$  element is mapped to.
- type(op\_map) :: *map1*: in case of directly accessed op\_dat this value must be set to the OP2 reserved name **OP\_ID**. In case of global data, this value must be set to the OP2 reserved name **OP\_GBL**. In case of an indirectly accessed op\_dat, this must be the reference to an op\_map variable mapping the iteration set to the set over which the op\_dat is defined.
- op\_access :: *access1*: this parameter can assume the values *OP\_READ*, *OP\_WRITE*, *OP\_RW* or *OP\_INC* (see Overview Section).

```

subroutine userKernel ( in1, in2, ..., inN )

    ! currently supported data types are integer(4), real(4) and real(8)
    "type1", dimension(size1) :: in1
    "type2", dimension(size2) :: in2
    ...
    "typeN", dimension(sizeN) :: inN

    ! use of input variables

end subroutine userKernel

```

Figure 5: General form of a user kernel.

## 7 User Kernel Interface

In this section we show how to express a user kernel to be used in a `op_par_loop`. The general form of a user kernel is shown in Figure 5.

Each input data corresponds to a `op_dat` argument line in the caller `op_par_loop`. As above, “type” can be an arbitrary data type. The size of the passed input data corresponds to the *dimension* parameter passed to the *op\_decl\_dat* call when declaring the corresponding `op_dat` variable.

A full example including the declaration of `op_dat` variables, the invocation of an `op_par_loop` and the corresponding user kernel declaration is shown in Figures 6 and 7.

```

program main

  real(8), dimension(:), allocatable :: x, q, adt, res

  ! allocation and filling of x, q, adt and res, declaration of names, etc..

  ! name example
  character(len=5) :: itSetName = 'nodes'

  ! declaration of iteration set
  type(op_set) :: iterationSet

  ! declaration of another set
  type(op_set) :: anotherSet

  ! mapping between iterationSet and anotherSet
  type(op_map) :: iter2another

  ! declaration of op_dat variables
  type(op_dat) :: p_x, p_q, p_adt, p_res

  ! initialisation of sets
  call op_decl_set ( itSetSize, iterationSet, itSetName )
  call op_decl_set ( anothSetSize, anotherSet, anothSetName )

  ! initialisation of map
  call op_decl_map ( iterationSet, anotherSet, 2, mappingData, iter2another,
    mapName)

  ! initialisation of op_dat variables
  call op_decl_dat ( iterationSet, xDimension, x, p_x, xName )
  call op_decl_dat ( anotherSet, qDimension, q, p_q, qName )
  call op_decl_dat ( anotherSet, adtDimension, adt, p_adt, adtName )
  call op_decl_dat ( anotherSet, resDimension, res, p_res, resName )

  ! parallel computation
  op_par_loop ( res_calc, iterationSet, &
    & p_x, -1, OP_ID, OP_READ, &
    & p_q, 1, iter2Another, OP_READ, &
    & p_q, 2, iter2Another, OP_READ, &
    & p_adt, 1, iter2Another, OP_READ, &
    & p_adt, 2, iter2Another, OP_READ, &
    & p_res, 1, iter2Another, OP_INC, &
    & p_res, 2, iter2Another, OP_INC &
    & )

end program main

```

Figure 6: Full example to show how a user kernel is defined. In this figure we show the code of the main program calling the `op_par_loop`.



```

subroutine res_calc ( x1, q1, q2, adt1, adt2, res1, res2 )

    ! formal parameters
    real(8), dimension(2) :: x1
    real(8), dimension(4) :: q1
    real(8), dimension(4) :: q2
    real(8), dimension(1) :: adt1
    real(8), dimension(1) :: adt2
    real(8), dimension(4) :: res1
    real(8), dimension(4) :: res2

    res1(1) = res1(1) + 0.5 * ( x1(1) + q1(1) + adt1(1) )
    res1(2) = res1(2) + 0.5 * ( x1(2) + q1(2) + adt1(1) )
    res1(3) = res1(3) + 0.5 * ( x1(1) + q1(3) + adt1(1) )
    res1(4) = res1(4) + 0.5 * ( x1(2) + q1(4) + adt1(1) )

    res2(1) = res2(1) + 0.5 * ( x1(1) + q2(1) + adt2(1) )
    res2(2) = res2(2) + 0.5 * ( x1(2) + q2(2) + adt2(1) )
    res2(3) = res2(3) + 0.5 * ( x1(1) + q2(3) + adt2(1) )
    res2(4) = res2(4) + 0.5 * ( x1(2) + q2(4) + adt2(1) )

end subroutine res_calc

```

Figure 7: Full example to show how a user kernel is defined. In this figure we show the code of the user kernel.